



Exploratory inference costs: four RTX 3090 cards, BF16, TP4, concurrency 1

Warmed latency distribution



Generated output usage



No project download or package installation ran during these measurements. Engine differences can confound latency. Native thinking defaults may differ; tokens from different tokenizers are not equivalent units of work.